



# SAP on Azure DevOps for SAP Workloads

Wednesday, October 21st, 2020  
10am – 12 pm SGT

Nicolas Yuen & Ravi Gangampali  
Microsoft APAC

**IMPT NOTICE:**

- If you choose to participate in this session using Microsoft Teams, your name, email address, phone number, and/or title may be viewable by other session participants.
- Please note that the training will not and cannot be recorded in alignment with Microsoft's policies



# SAP on Azure Partner Enablement

Module Two – Week Two

Day 3 – SAP and DevOps



**Nicolas Yuen**  
Cloud Solution Architect



**Ravi Gangampalli**  
Cloud Solution Architect – SAP on Azure



**Anjan Banerjee**  
Sr Customer Engineer – SAP on Azure  
Azure Engineering

# Check-in

We are happy to host you 😊

<https://aka.ms/apac-enablement-check-in>



---

# Agenda

DevOps – quick recap  
SAP deployment on Azure  
SAP and Terraform  
Infra As Code core concepts  
Automation with Azure DevOps

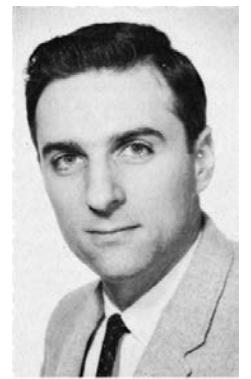
# DevOps quick Recap

# What is DevOps

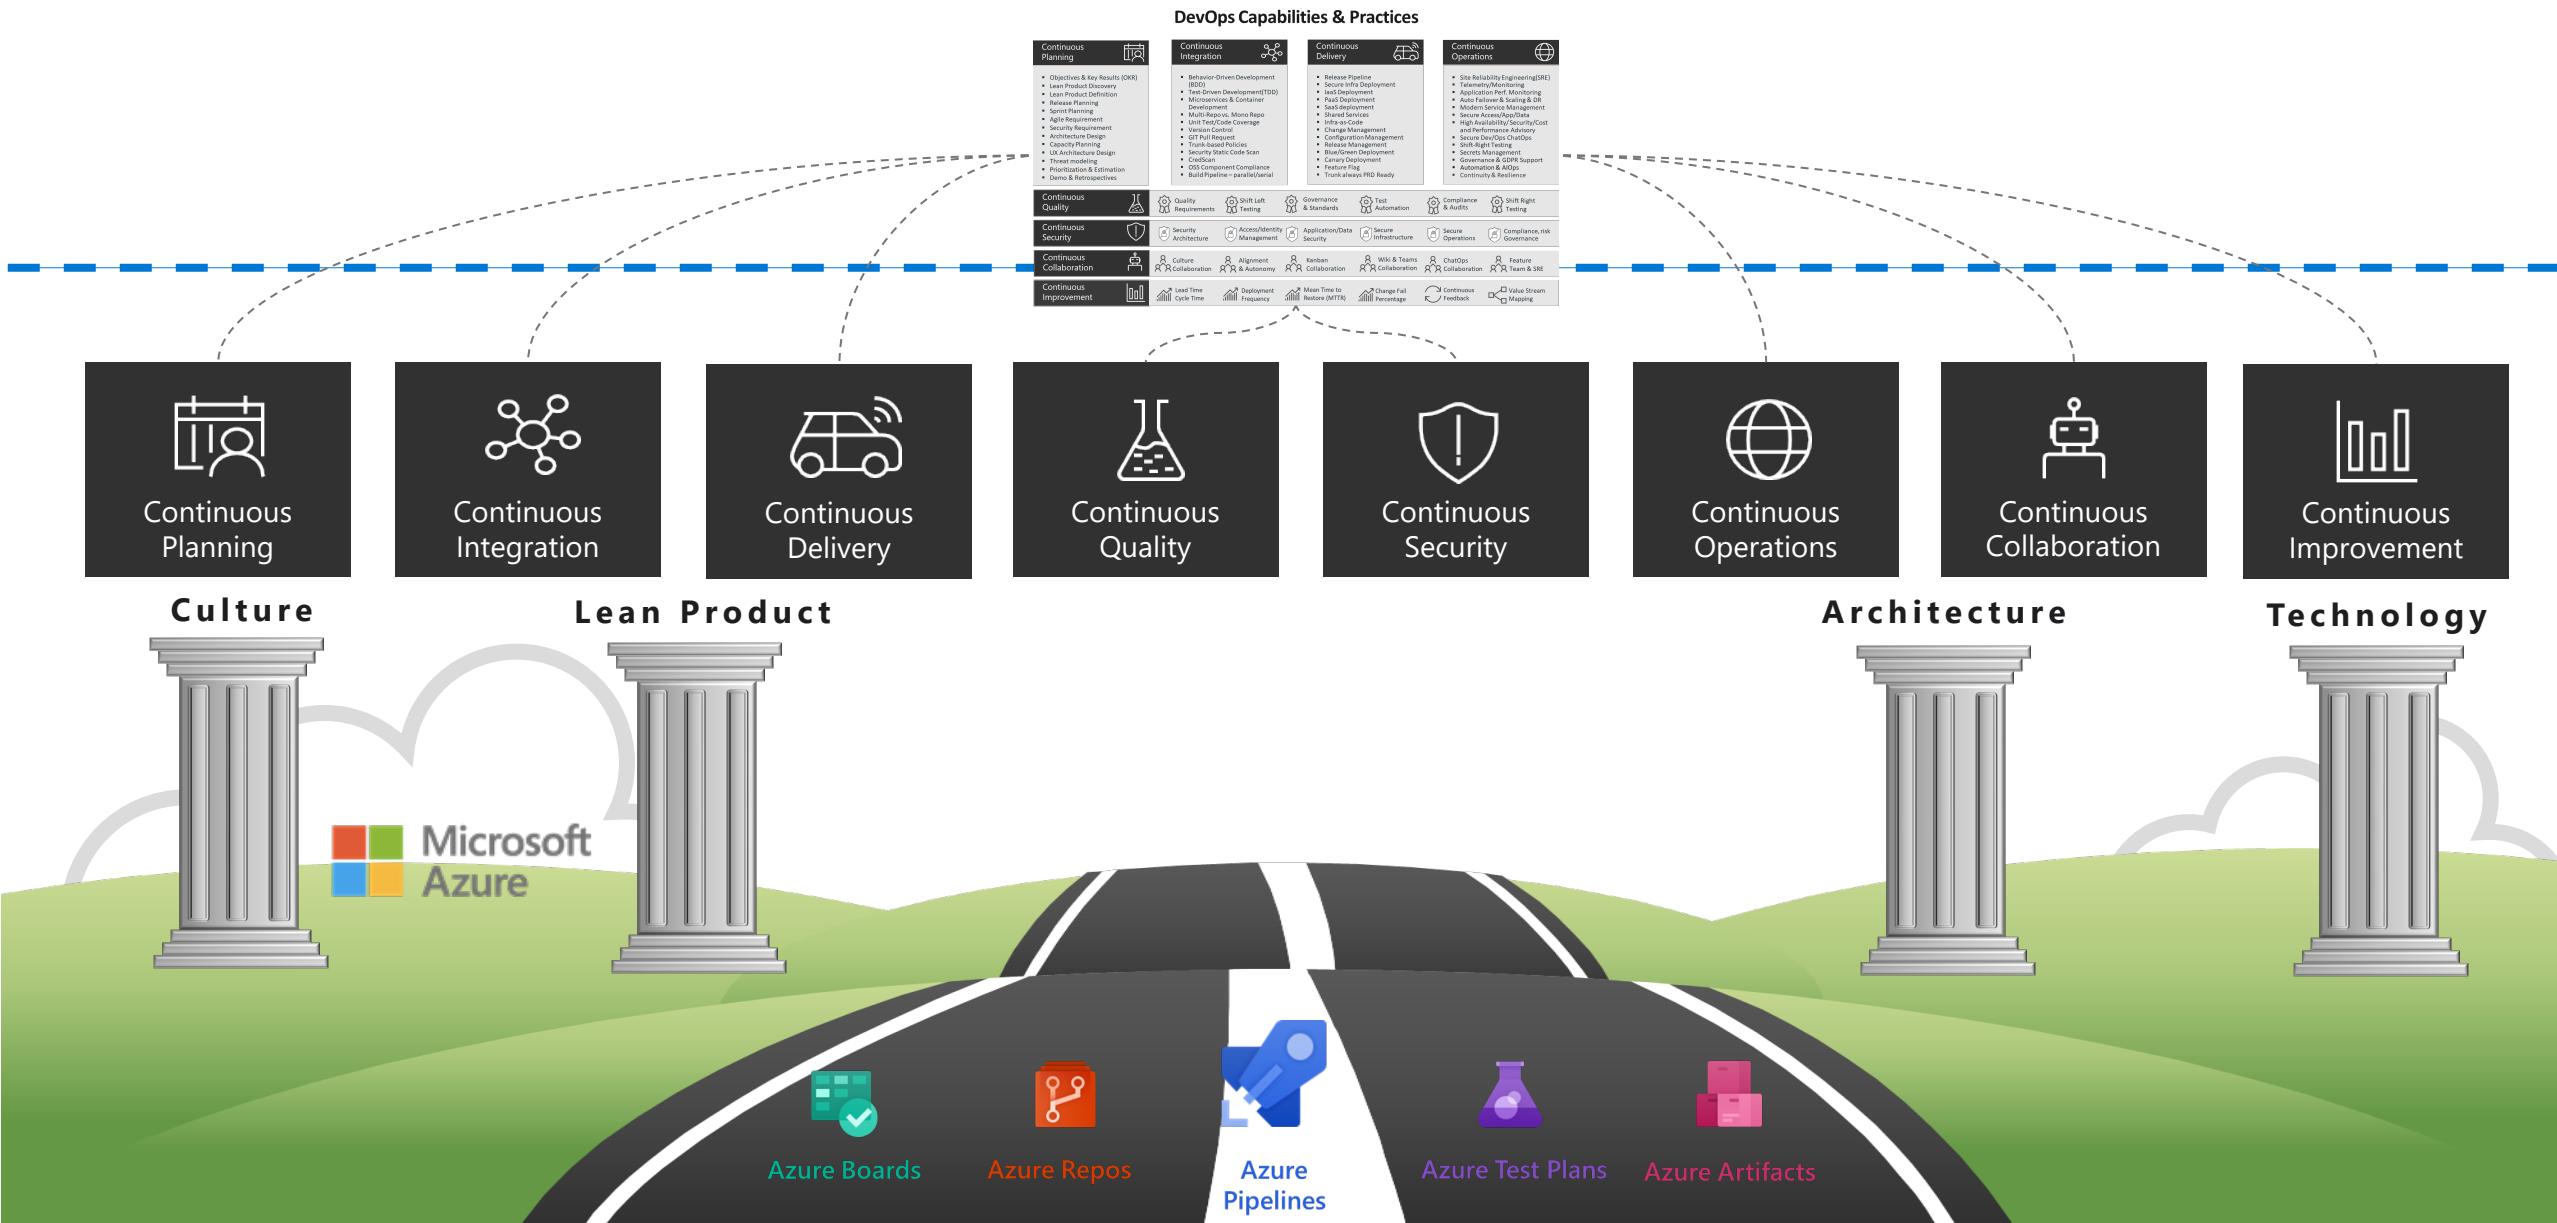
The contraction of “Dev” and “Ops” refers to **replacing siloed** Development and Operations to **create multidisciplinary teams** that work together with **shared and efficient practices and tools**.

# Conway's law

- “Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure” Melvin Conway – 1968
- “*The org chart of the company has a direct impact on the code*” Clément Rochas



# DevOps Taxonomy



# Mars Climate Orbiter

---

- **Mission**

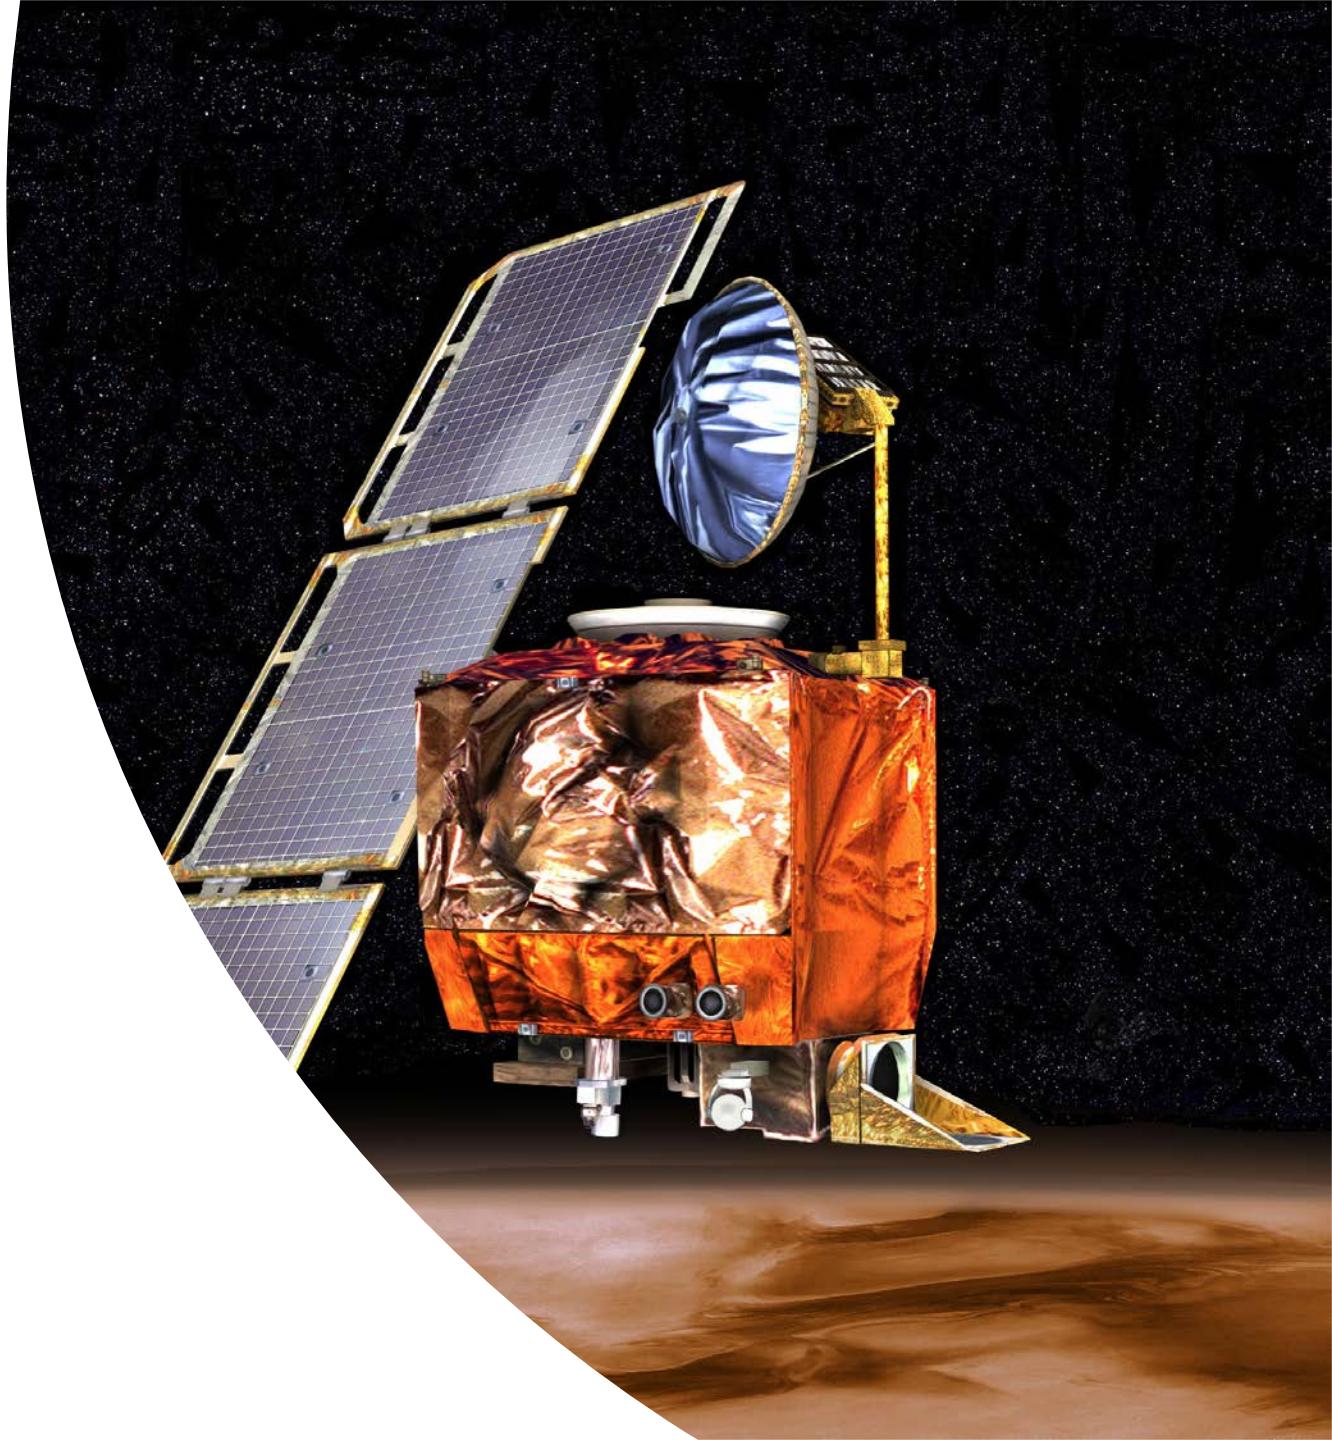
The craft was intended to be the first ever weather satellite for Mars, which would monitor the planet's atmosphere with its high-resolution camera.

- **What happened**

- The craft burned up in the Martian atmosphere
- \$327.6 million loss for NASA

- **What went wrong**

- Bug in the ground control software supplied by a third party
- Calculations were performed manually rather than using the supplied software, due to file format errors and miscellaneous bugs



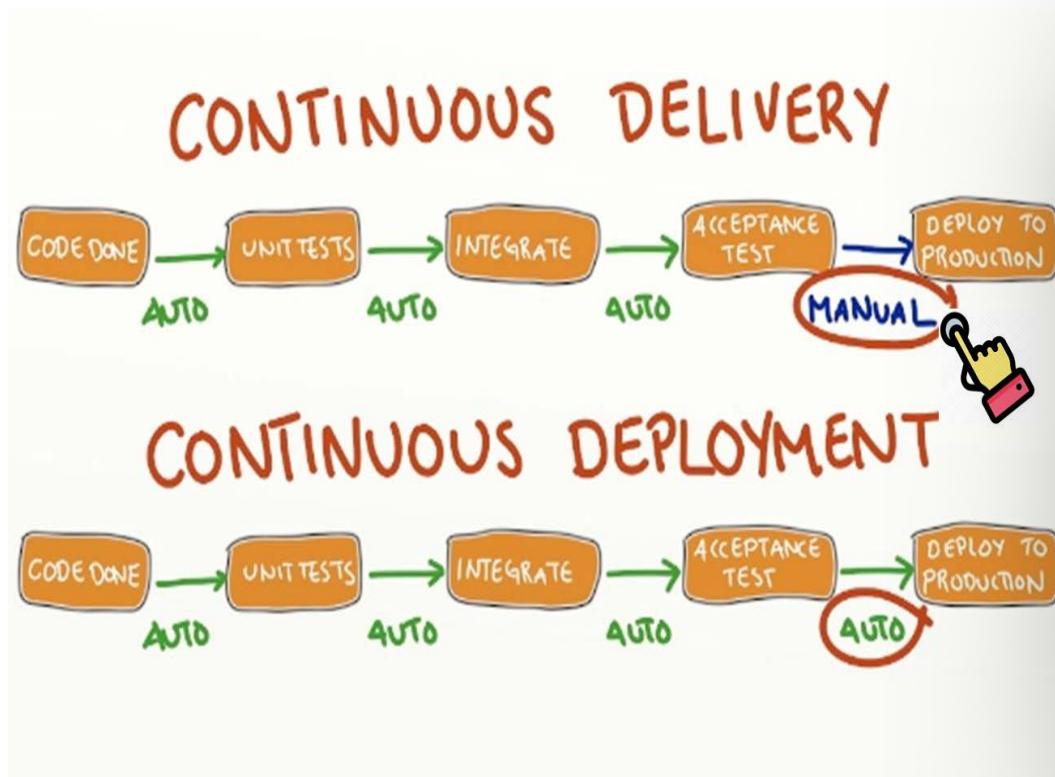
# What are the Goals of Continuous Integration

The main purpose of Continuous Integration is to prevent developers stepping over each other code and eliminate integration issues.

## **CI Goals:**

- #1: Harness collaboration**
- #2: Enable parallel development**
- #3: Minimize integration debt**
- #4: Act as a quality gate**
- #5: Automate Everything!**

# What is Continuous Delivery



**Continuous delivery** is a software engineering approach in which teams produce software in short cycles, ensuring that the software can be reliably released at any time and, when releasing the software, doing so manually. It aims at building, testing, and releasing software with greater speed and frequency. The approach helps reduce the cost, time, and risk of delivering changes by allowing for more incremental updates to applications in production. A straightforward and repeatable deployment process is important for continuous delivery.

# Continuous Integration Practices



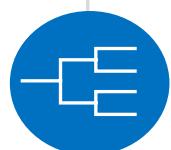
Multi-Repo



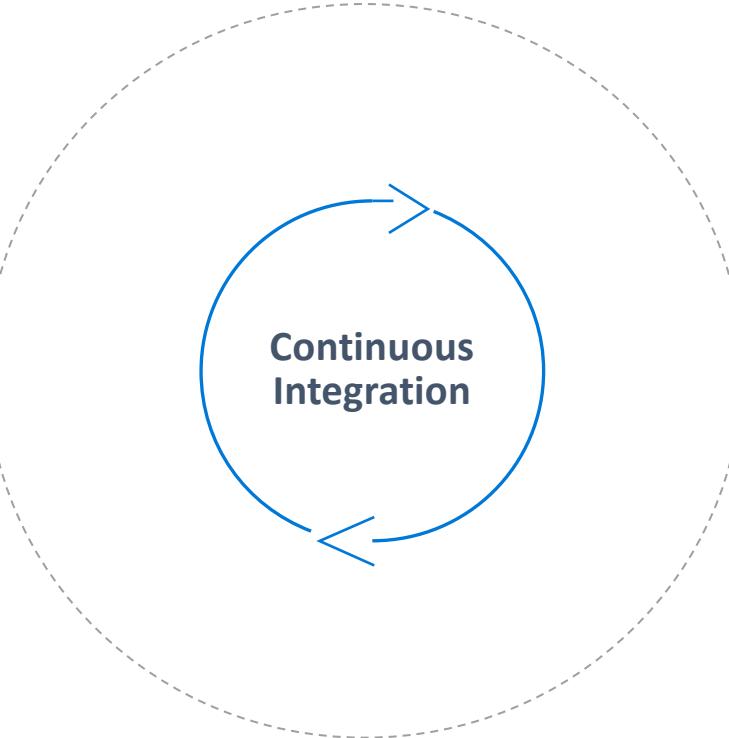
Trunk Based Branching



Branch Policies



Parallel Builds



Culture & Collaboration



Unit Testing



Pull Requests



Security Scanning



# Continuous Delivery Practices



Automated Builds



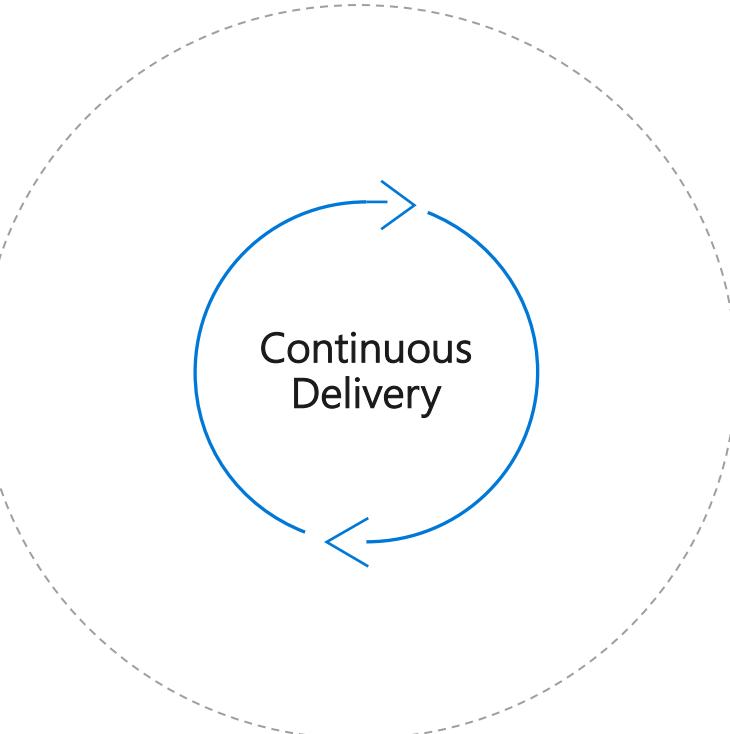
Build Once



Trunk Always Prod  
Ready



Fix Immediately



Everything in Source Control



Feature Flags



Blue/Green Deploy



Canary Deployment



# SAP deployment options

# Methods for deploying SAP (and everything else) in Azure

Three over-arching methods.

- Azure Portal (manual)
- Scripts / SDKs (automation)
- Template based deployments (automation)

# Methods for Azure resource deployments

## Azure Portal (manual)

Pros:

- Browser based, quick setup, no fuss
- Nice for exploration and visual inspection & monitoring
- Fully featured

Cons:

- Everything is performed manually
- Error prone
- Lack of process integration (DevOps, ITSM)

# Methods for Azure resource deployments

## Scripts / SDKs (automation)

Pros:

- Process integration (DevOps / ITSM) possible
- Removes human / less error prone
- Unopinionated / total flexibility

Cons:

- Requires scripting knowledge / environment
- Complex logic needs to be hand built
- DIY error handling

# Creating a RG, vnet, & VM with Scripts

- Powershell Example here: <https://docs.microsoft.com/en-us/azure/virtual-machines/scripts/virtual-machines-windows-powershell-sample-create-vm>
- CLI example here: <https://docs.microsoft.com/en-us/azure/virtual-machines/scripts/virtual-machines-windows-cli-sample-create-vm>

# Methods for Azure resource deployments

## Template based deployments (automation)

Pros:

- Process integration (DevOps / ITSM) possible
- Removes human / less error prone
- Handles some complex logic
- Options for state management

Cons:

- Requires templating knowledge / environment
- Opinionated and lack of full flexibility
- Some error handling

# Template based deployments

Digging deeper on template based deployments.

- Azure Resource Manager templates or Terraform
- Declaration of desired infrastructure
- JSON or JSON like syntax
- Deploy, update, delete

# Terraform

## What is Terraform?

- Open source project
- Cross computing environment templating language
- Provision, Update, and Delete resources
- Authored in HashiCorp Configuration Language (HCL) or JSON

# Terraform Example

```
resource "azurerm_resource_group" "testrg" {  
    name = "resourceGroupName"  
    location = "westus"  
}
```

```
resource "azurerm_storage_account" "tests" {  
    name = "storageaccountname"  
    resource_group_name = "testrg"  
    location = "westus"  
    account_tier = "Standard"  
    account_replication_type = "GRS"  
}
```



Resource Group

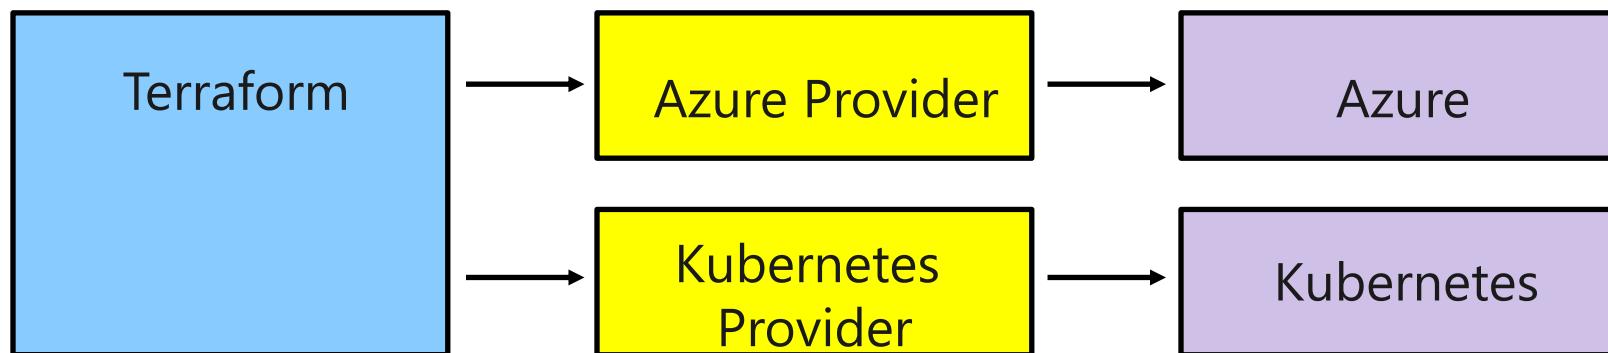


Storage Account

# Providers

What is a Terraform provider?

- Terraform 'extensions' for deploying resources
- Manages cloud / endpoint specific API interactions
- Available for major clouds and other platforms
- Hand authored (azurerm)



# Basic resource creation

## Deployment foundations.

- Resource Type: required provider
- Name: internal name
- Configuration: deployment details

```
Resource Type          Name  
resource "azurerm_resource_group" "demo-rg" {  
    name = "demo-rg"  
    location = "westus" } Resource Configuration
```

# Automation, orchestration & Configuration Management Tools

Once we have created the infrastructure, we need to do something with it

- Install Corporate standards
- Establish Authentication / Domain membership
- Install SAP prerequisites
- Install SAP software
- Install other supporting software
- Maintain ongoing configuration
- Perform patches, updates, etc

# Automation, orchestration & Configuration Management Tools

Wide Variety of tools available:

- Custom Script Extension & shell script
- Puppet, Chef, Ansible
- Azure DSC
- Etc.

# Checklist

Resource group has been setup appropriately	SAP License – SLICENSE
RBAC access for each resource group is setup	Initial consistency check - SICK
System deployed in appropriate region	Application instances – SM51
Ensure resource locks are setup	Operation mode check – RZ03
Compute size chosen for SAP application is per certified per SAP note 1928533	Web Dispatcher
Right choice of OS is made for M Series (SUSE 12 SP3 or RHEL 7.4)	Transport Management System
Storage configured in the VM and striped appropriately	SMQ1/SMQ2
Read cache enabled for database data disks	email (SCOT)
Write accelerator enabled for HANA log disks	Check LC10, DB59, and connect to Live Cache from SAP
Accelerated networking is enabled	Check all the important RFC destinations and critical interfaces
Network security group and firewall is setup appropriately (consider RFC connections)	Enqueue Replication incase of HA
Subscription management server is setup	SMD agents are installed
OS patches are installed and upto date	Ensure connectivity to solution manager
sapline parameters is applied for HANA DB VM's	DBACOCKPIT: Main database analysis tasks, SQL editor needs to be available
OS parameters per SAP note is configured (parameters differs for every release and service pack of operating system)	AL11: Directory access
Ensure ER gateway is configured	RSPFPAR: SAP profile parameters
Ensure SAP systems are deployed in the appropriate VNETs	SE11: ABAP dictionary
configure VNET peering if required	SE37: Execution of function modules (DB_EXECUTE_SQL, ...)
Ensure cluster parameters are in place	SE38 / SA38: Execution of reports
Perform failover testing after cluster configuration	SM21: SAP syslog (logs from all instances)
Ensure appropriate I/O fencing mechanisms are configured	SM49: Execution of logical commands
Ensure NTP is setup	SM50: Instance specific SAP work process overview
Time is synchronized across application and databases VM's	SM51: SAP instance overview
Ensure appropriate mountpoints are setup	SM66: Global SAP work process overview
for SAP HANA parameters refer to SAP Note 1969700 - SQL Statement Collection for SAP HANA	SNRO: Number range information
Refer SAP Note: 2186744 - FAQ: SAP HANA Parameters	ST02: SAP buffers and memory
	ST03 / ST03N: SAP workload analysis
	ST05: Execution plans, SQL trace
	ST06: Operating system analysis
	ST10: SAP table buffer analysis
	ST22: SAP short dumps
	STAD: SAP statistical records

# Infrastructure as Code + key concepts

# Infrastructure as code



Stand up environments in the fastest means possible.



Reduce human errors and increase reliability.



Improve environment visibility and improve developer efficiency



Store your infrastructure definitions alongside your application code.

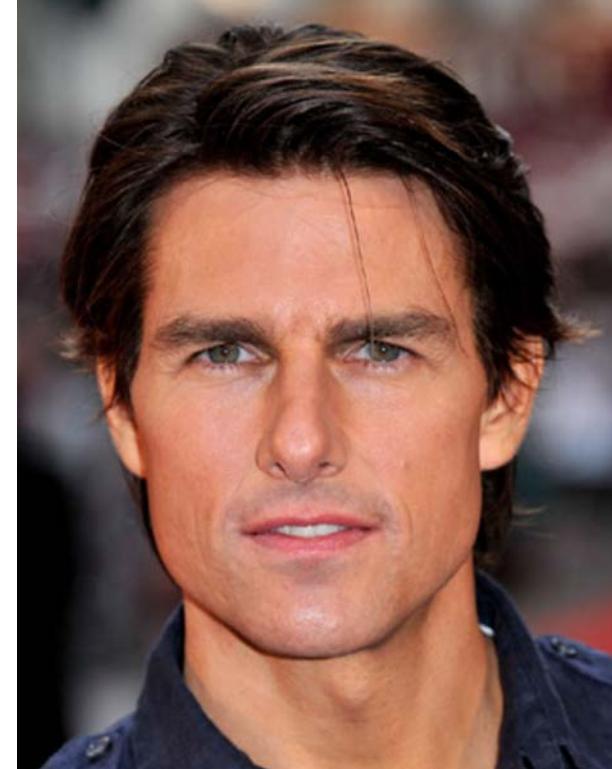
```
File Edit Selection View Go Debug Terminal Help aks.tf - terraform - Visual Studio Code

EXPLORER OPEN EDITORS
  ✓ aks.tf
  ✓ TERRAFORM
    > .terraform
      ✓ aks.tf
      ✓ azuread.tf
      ✓ gateway.tf
      ✓ helm.tf
      ✓ kubernetes.tf
      ✓ main.tf
      ✓ monitoring.tf
      ✓ networking.tf
      { terraform.tfstate
      terraform.tfstate.backup
      ✓ variables.tf

aks.tf
1 resource "azurerm_kubernetes_cluster" "default" {
2   name          = "${var.name}-aks"
3   location      = "${azurerm_resource_group.default.location}"
4   resource_group_name = "${azurerm_resource_group.default.name}"
5   dns_prefix    = "${var.name}-aks-${var.environment}"
6   depends_on    = ["azurerm_role_assignment.default"]
7   kubernetes_version = "1.14.0"
8
9   agent_pool_profile {
10     name        = "default"
11     count       = "${var.linux_node_count}"
12     vm_size     = "${var.linux_node_sku}"
13     os_type     = "Linux"
14     os_disk_size_gb = 30
15     vnet_subnet_id = "${azurerm_subnet.pod.id}"
16     type        = "VirtualMachineScaleSets"
17   }
18
19   agent_pool_profile {
20     name        = "win"
21     count       = "${var.windows_node_count}"
22     vm_size     = "${var.windows_node_sku}"
23     os_type     = "windows"
24     os_disk_size_gb = 30
25     vnet_subnet_id = "${azurerm_subnet.pod.id}"
26     type        = "VirtualMachineScaleSets"
27   }
28
29   service_principal {
```



# Pet vs Cattle

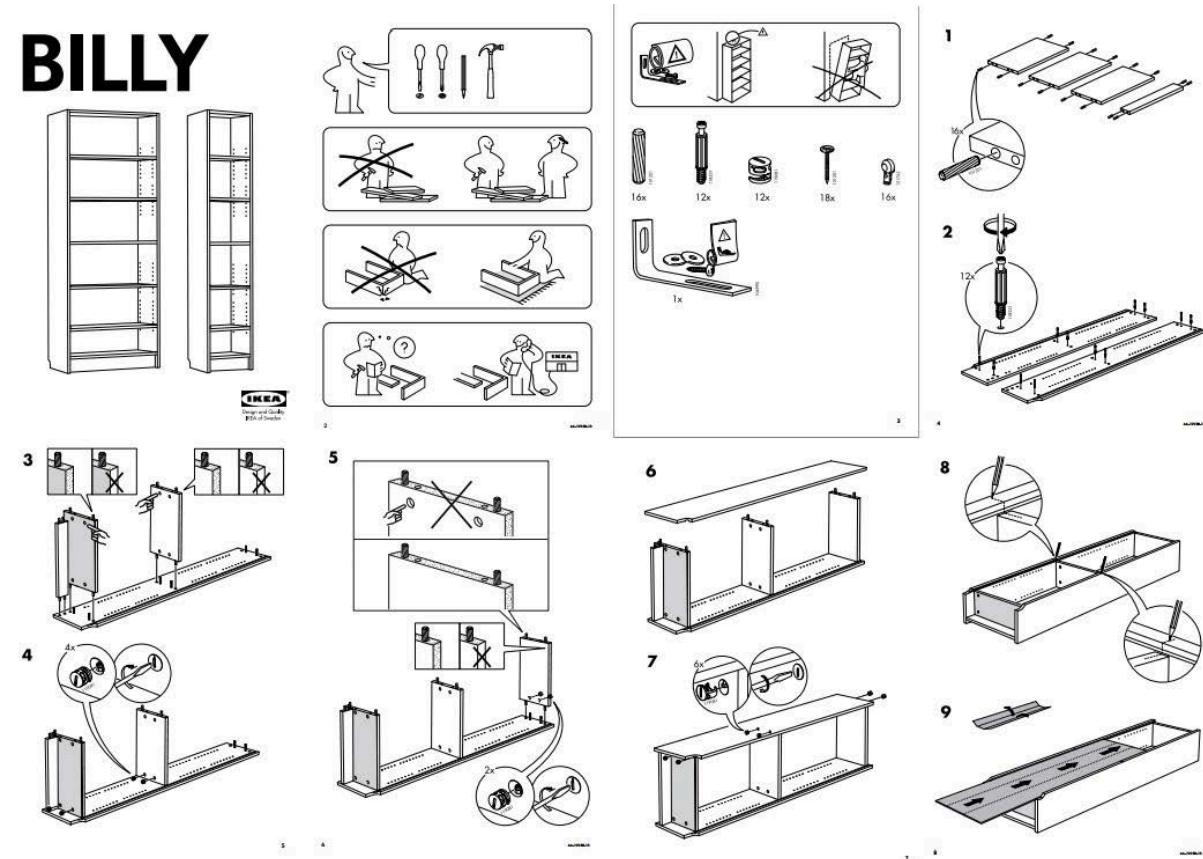


# Mutable vs Immutable



# Frying vs Baking

# BILLY



# Imperative vs Declarative



# Snowflake vs Phoenix

## Azure subscription

### GitOps

config  
registry

GIT  
repos

build  
agents

release  
agents

### Remote state management

storage  
account

VNET

keyvault

identity

### Launchpad

## Azure subscription with landing zones



Backup  
Simple and reliable server backup to the cloud



Azure Monitor  
Highly granular and real-time monitoring data for any Azure resource



Automation  
Simplify cloud management with process automation



Log Analytics  
Collect, search, and visualize machine data from on-premises and cloud

### Shared Services

Application  
infrastructure 1

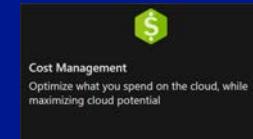
Application  
infrastructure n



Security & Compliance  
Enable threat detection and prevention through advanced cloud security



Azure Policy  
Implement corporate governance and standards at scale for Azure resources



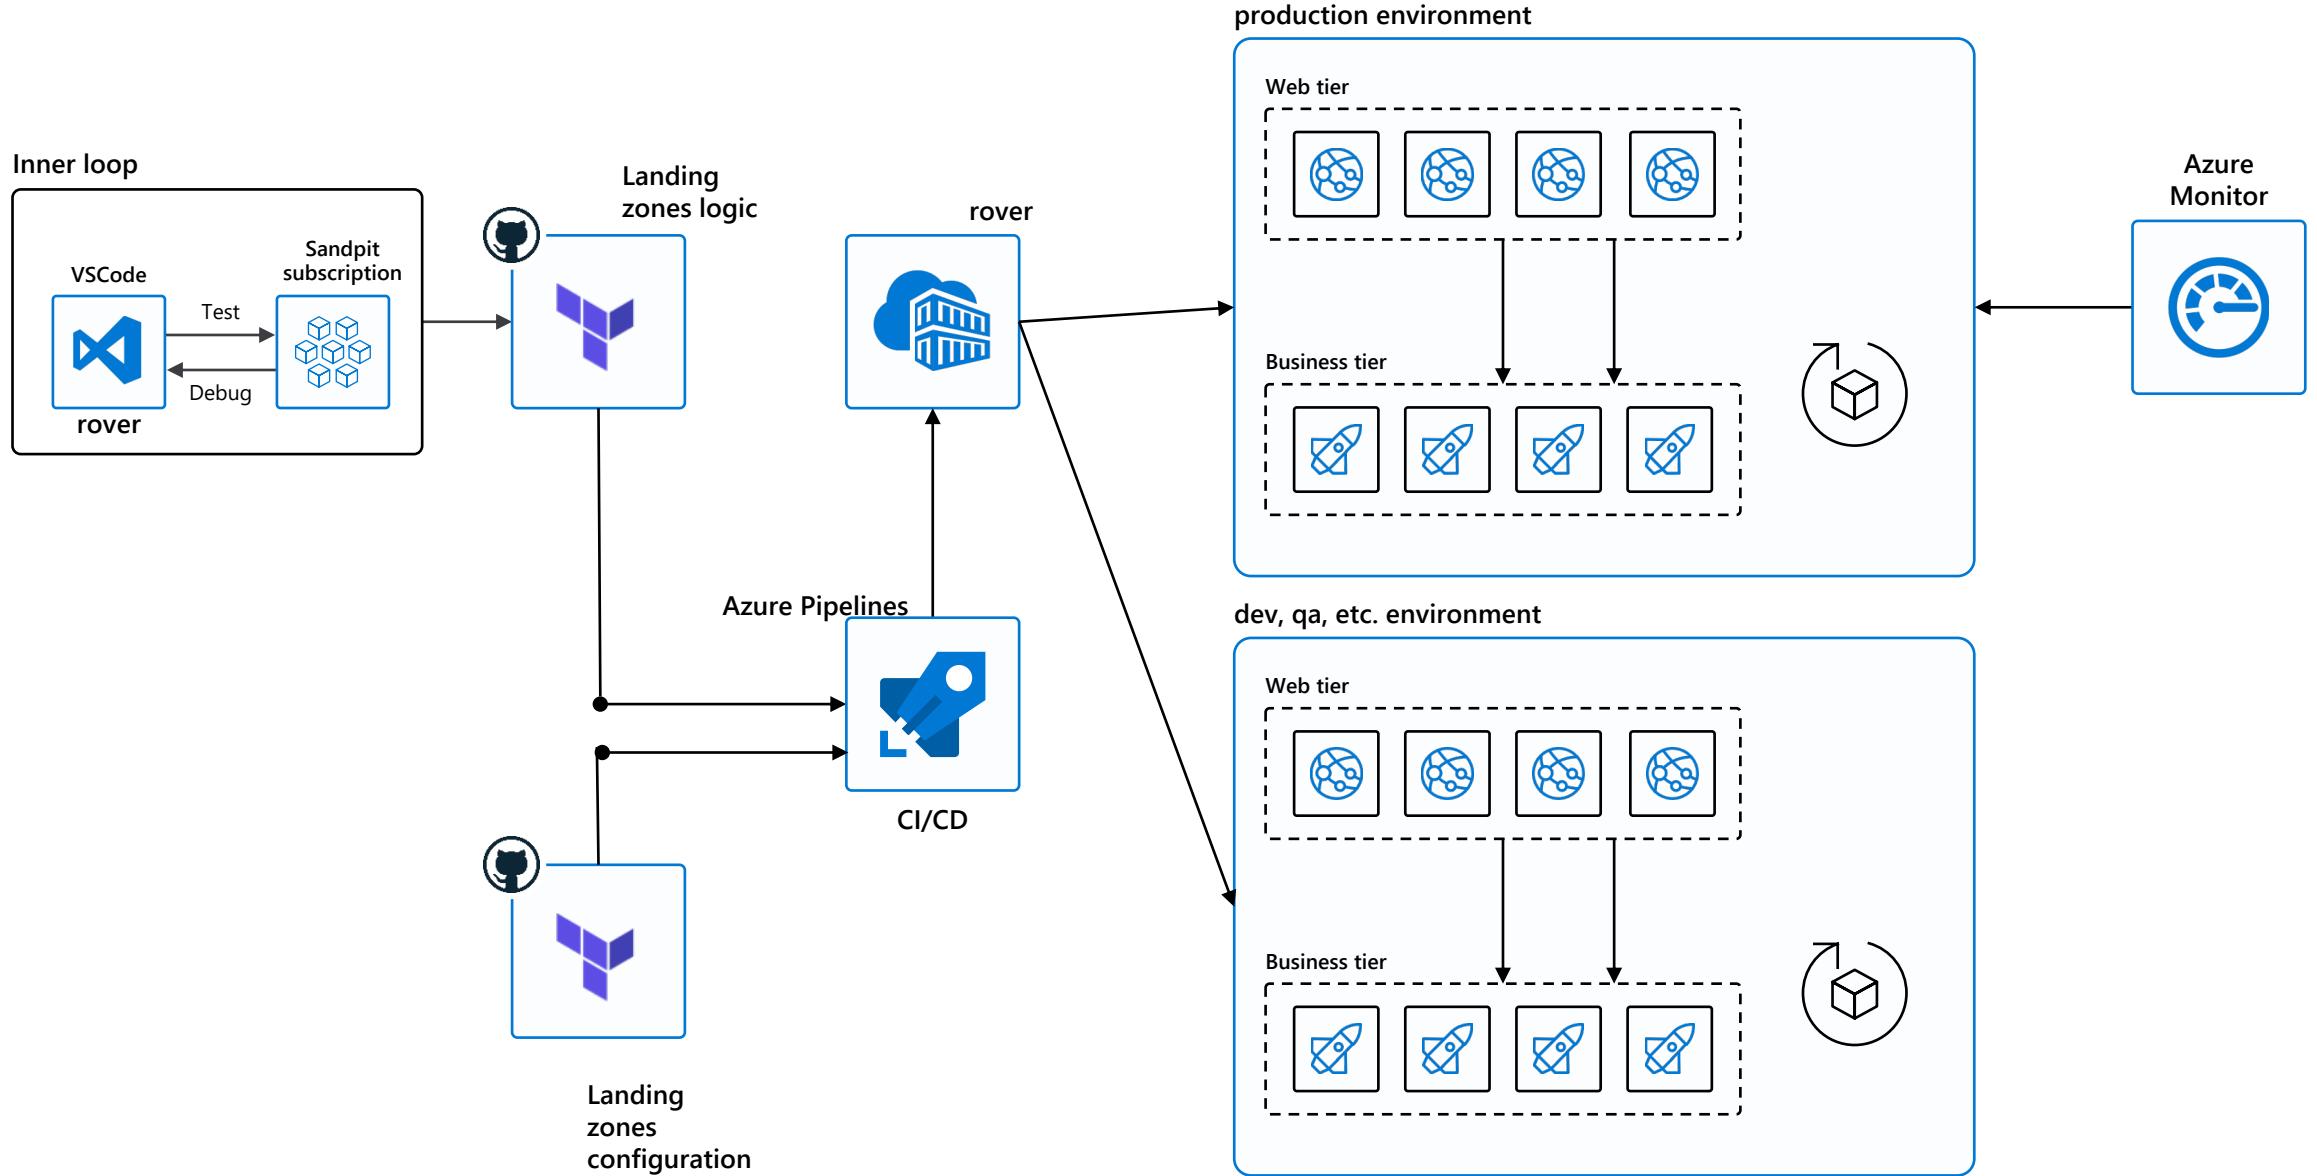
Cost Management  
Optimize what you spend on the cloud, while maximizing cloud potential

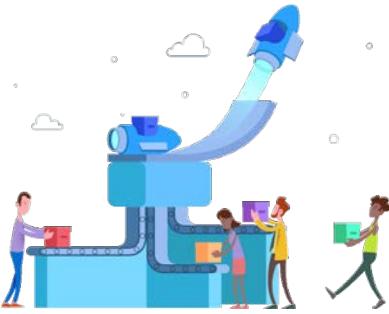


Azure Sentinel  
Learn how to use Microsoft's SIEM that provides intelligent security analytics for your entire enterprise at cloud scale

### Security and Governance

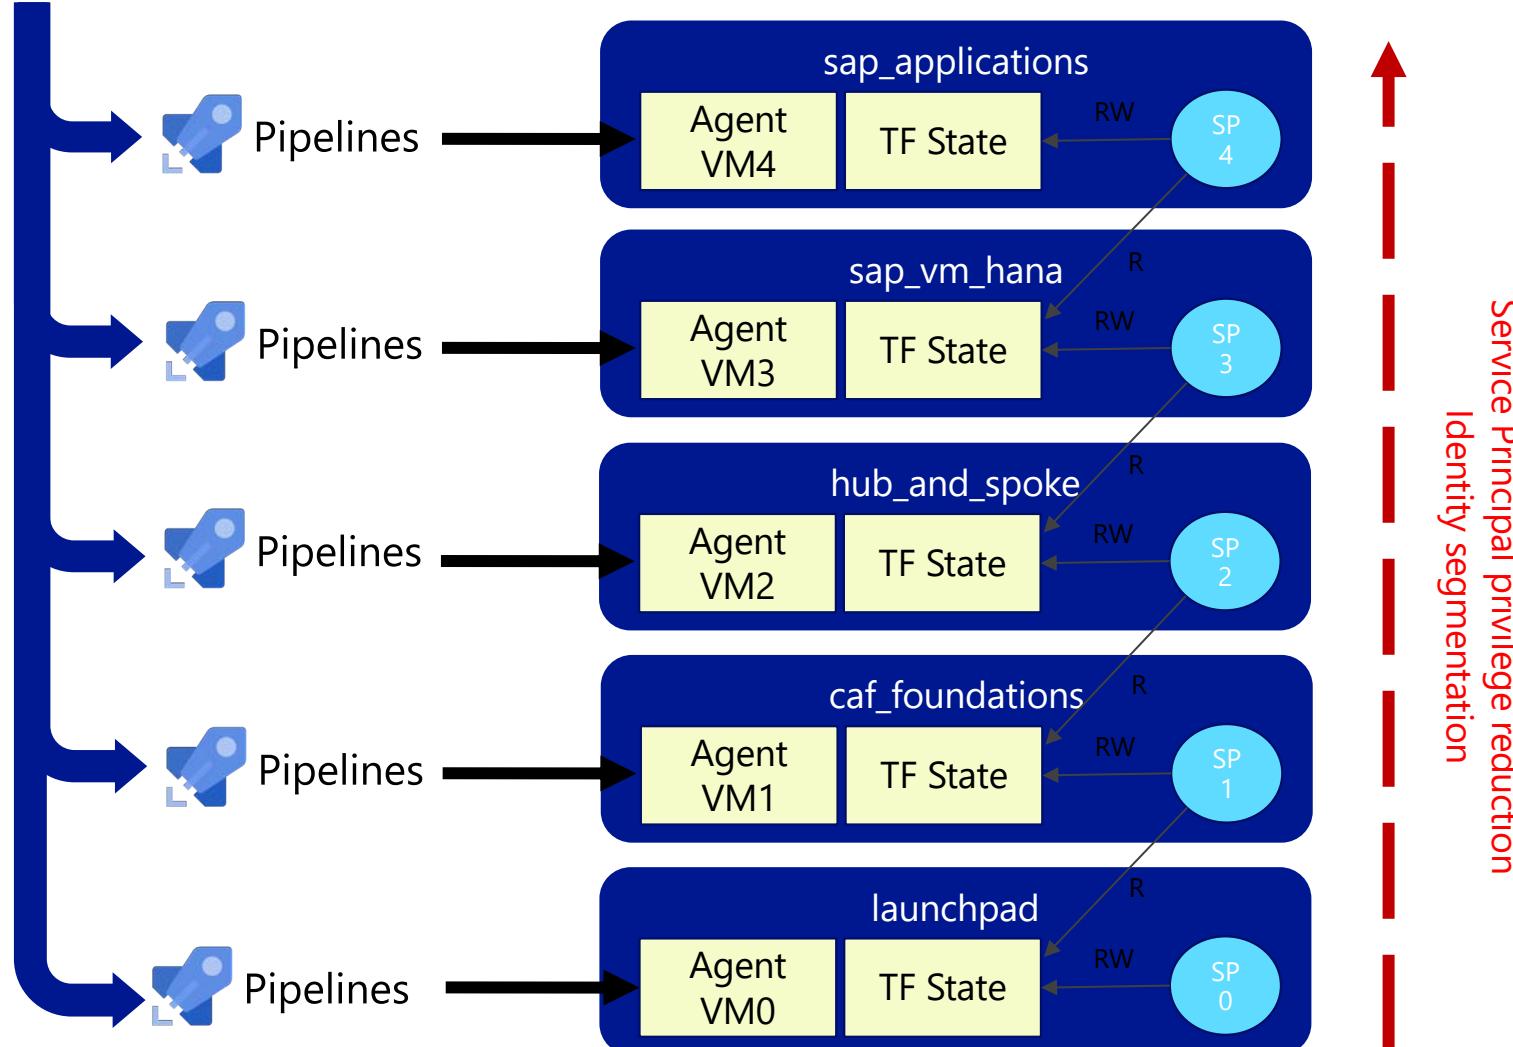
# Operating IaC - CI/CD Pipeline





## Azure Pipelines

Continuously build, test, and deploy to any platform and cloud



Landing zones structuring example

Manage the deployment of an **application itself / ML model** in the application's landing zone e.g (Springboot microservices, dotnet core...)

Manage the deployment of an **application's landing zone** in a spoke environment (e.g AKS Cluster + WAF)

Manage the **network hub and spokes** and **shared services** of each environments, as described in the design document (backup, DR, Azure monitor, patch management...)

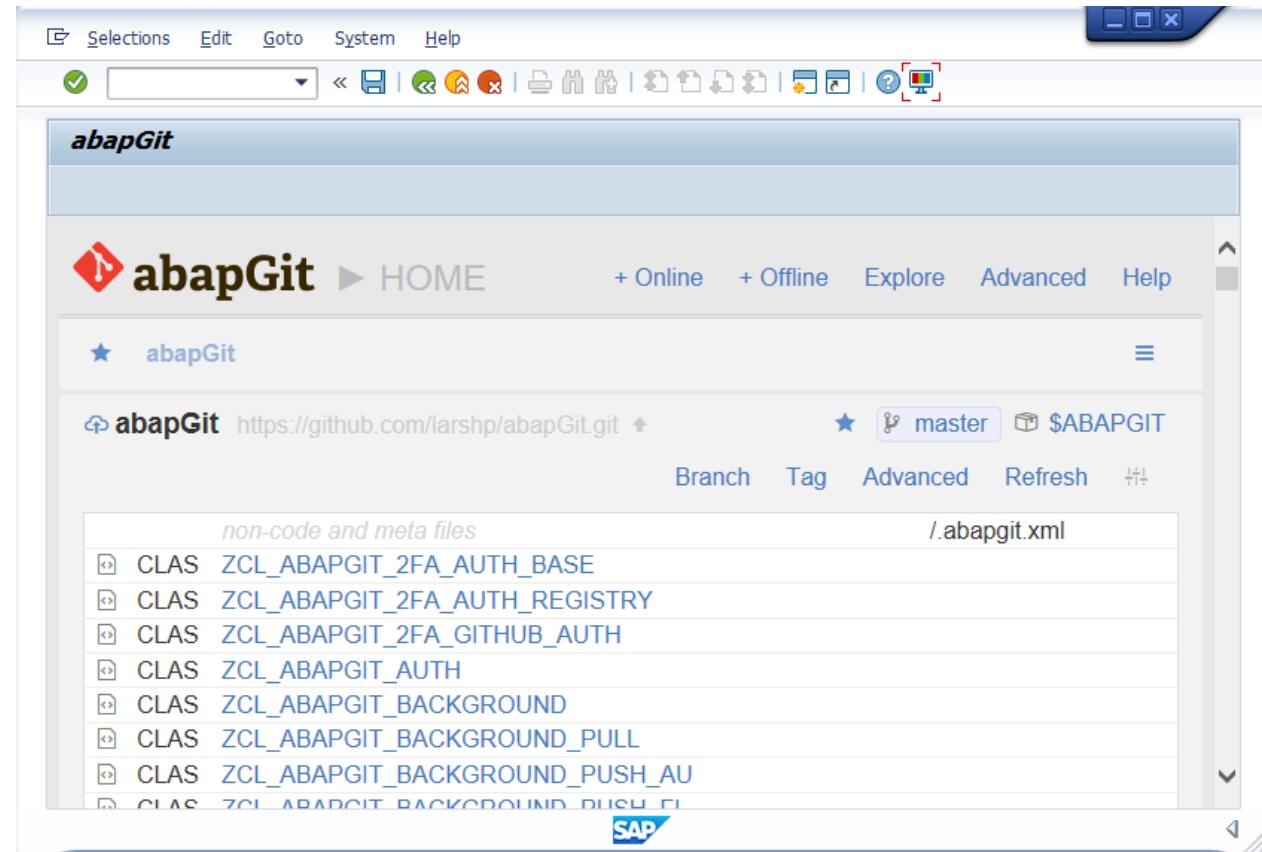
Manage the **security and compliance** (RBAC, Policies, OMS monitoring, shared security services, including event hub collector for SIEM, preventive and reactive controls...)

**Transition from manual to automation.** Create the subscriptions (for level 0 to level 4) + terraform state repository, privileged access workstation, service principals

# Automating SAP deployment

# abapGit – DevOps for SAP

- Created by [Lars Hvam](#) – Has many contributors
- Open Source project enabling SAP integration with Git
- Works on ABAP Package Level
- [User Exits allows Fine-Tuning](#)
- Single Sign-on using Azure Active Directory
- Integration with Azure DevOps



<https://github.com/Azure/sap-hana>



SAP Automation

## Automated SAP Deployments in Azure Cloud

Master Branch's status:  Azure Pipelines never built

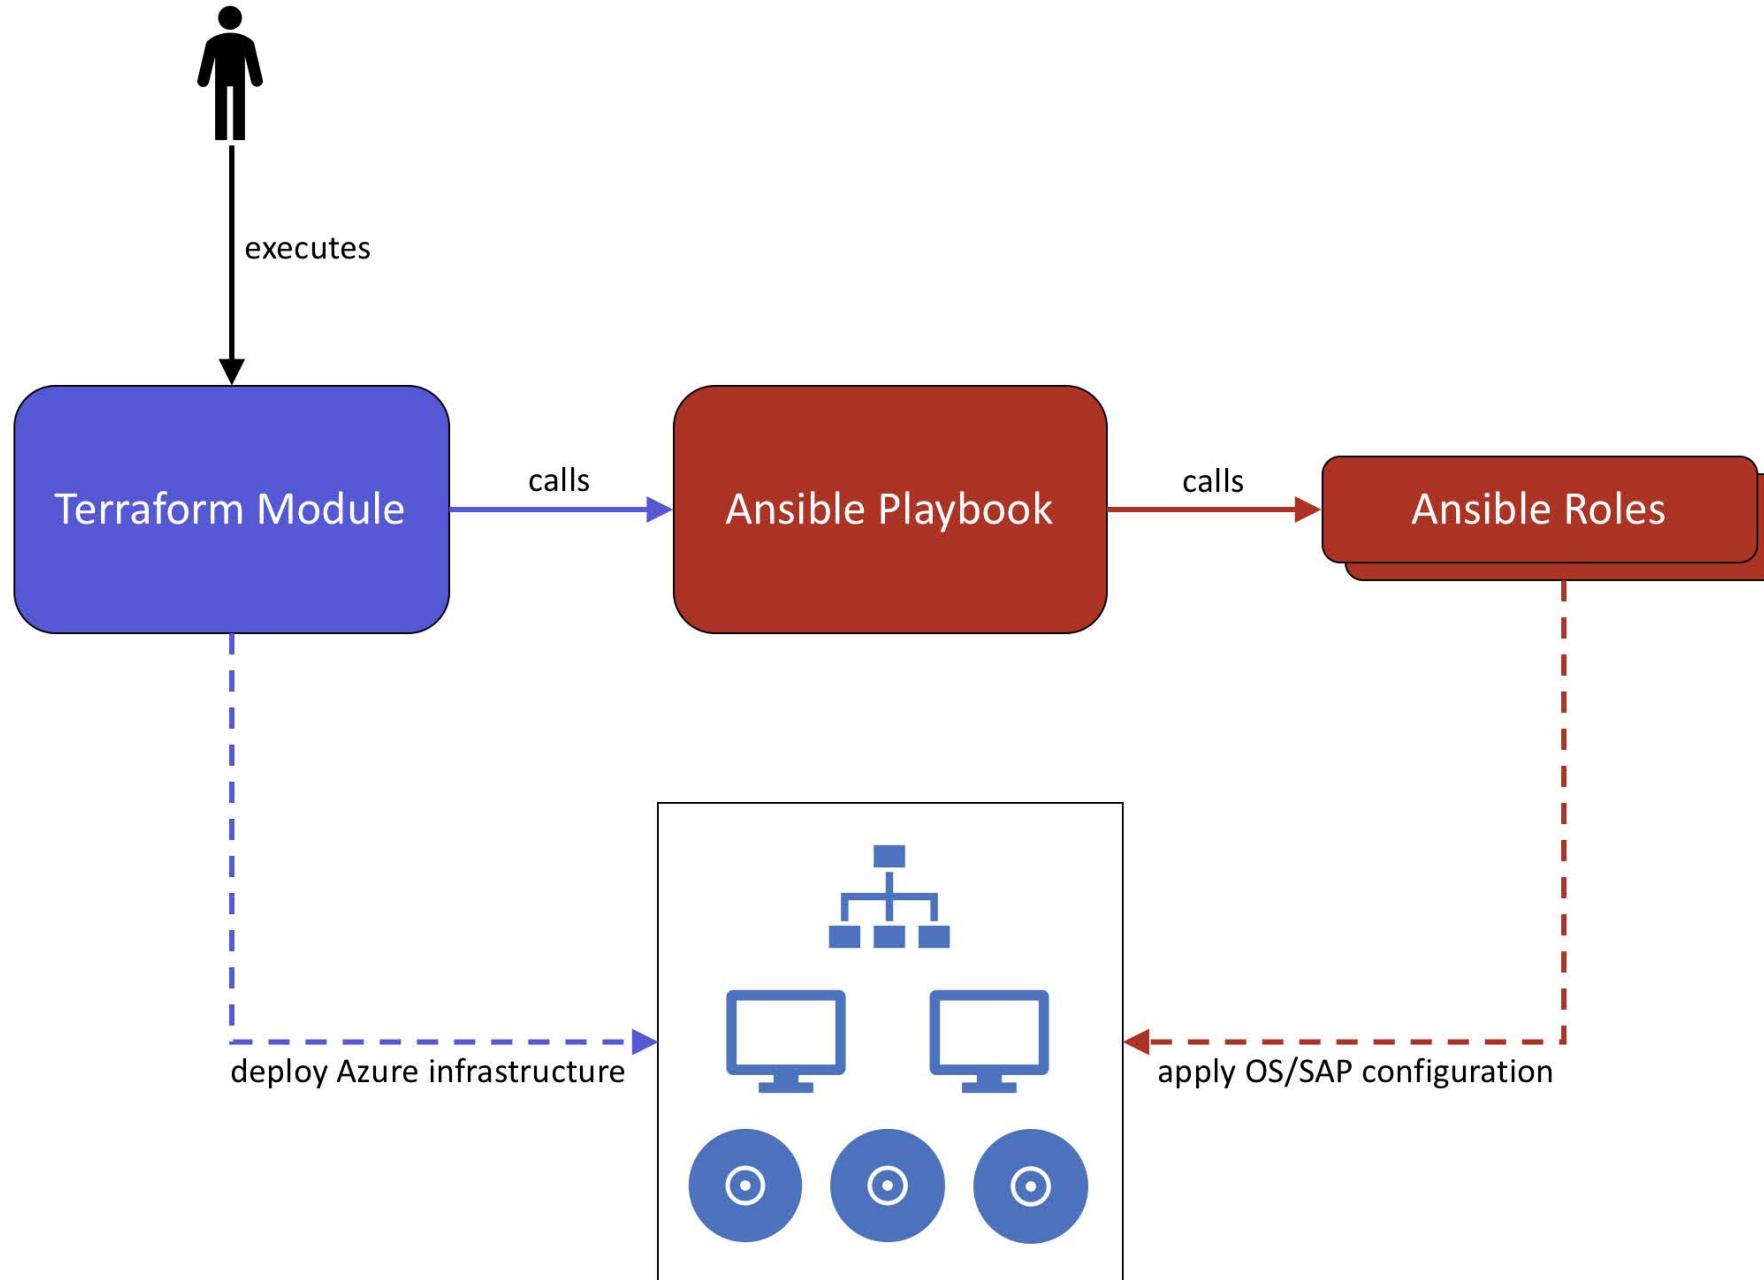
This repository contains code that can be used to automatically deploy SAP landscapes in the Azure Cloud.

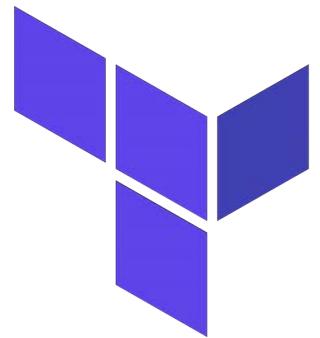
The templates are split into:

- **Terraform modules** which deploy the infrastructure components (such as VMs, network, storage) in Azure.
- **Ansible playbooks** which run different roles to configure and VMs and install SAP HANA and required applications on the already deployed infrastructure.

# A new approach to automating SAP

- Terraform modules which deploy the infrastructure components (such as VMs, network, storage) in Azure and then call the:
- Ansible playbook which call different:
- Ansible roles to install and configure OS and SAP applications on the deployed infrastructure in Azure.





# Supported Scenarios Available

- [HANA Scale-Up Stand Alone](#)
- [HANA with High Availability](#)

## Usage

- A typical deployment lifecycle will require the following steps:
  1. [Initialize the Deployment Workspace](#)
  2. [Adjusting the templates](#)
  3. [Running Terraform deployment](#)
  4. [Running Ansible Playbook](#)
  5. [Deleting the deployment](#) (optional)
  6. (*Note: There are some script under [sap-hana/util](#) would help if you are using Linux based workstation*)

# What will be deployed



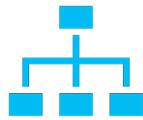
## Common Infrastructure

Resource Group  
Management VNet  
Management SubNet and Network Security Group (NSG)  
SAP VNet  
Storage Account for SAP Media downloads  
Proximity Placement Group  
iSCSI SubNet, NSG and Network Security Rules (NSR)  
iSCSI VMs



## HANA Database

Administration SubNet, NSG, and NSR  
Database SubNet, NSG, and NSR  
HDB Availability Set and Load Balancer  
HDB VMs with Data Disks



## Application Tier

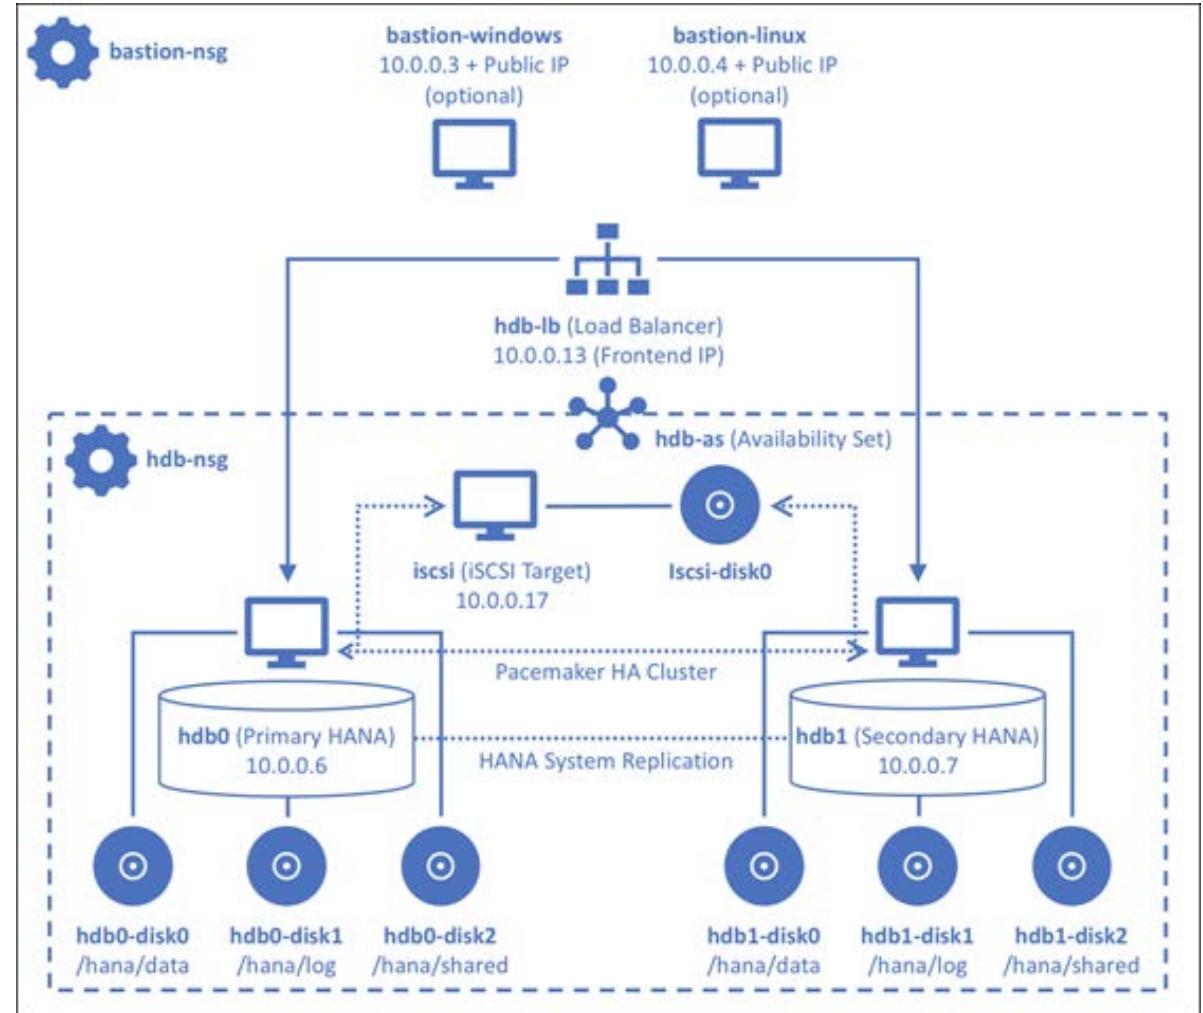
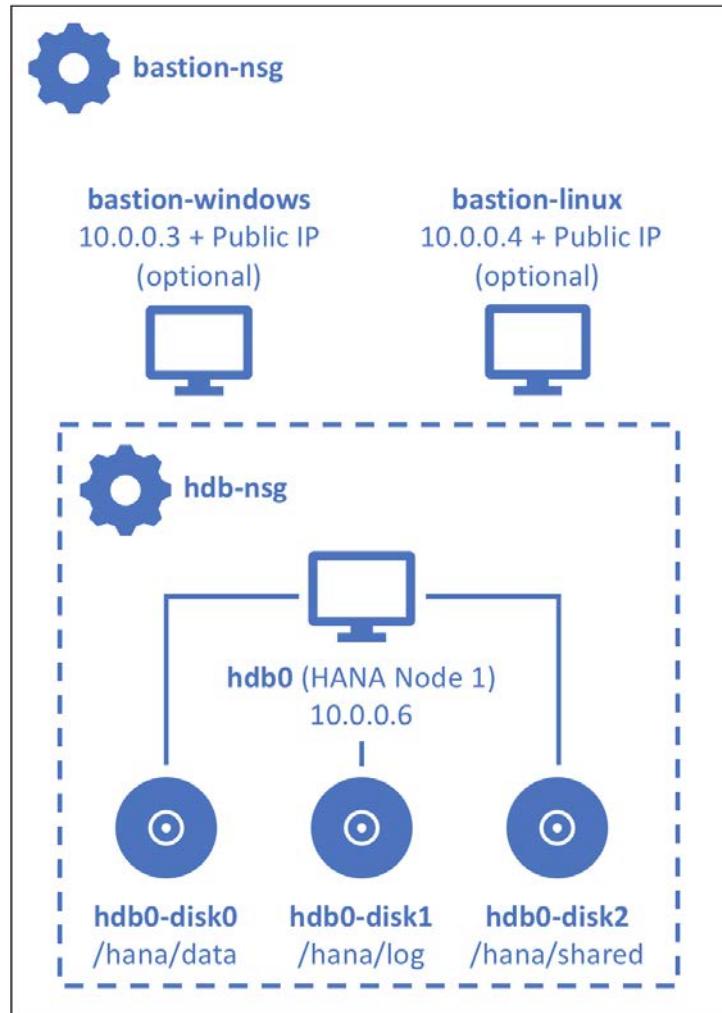
Application Tier SubNet, NSG, and NSR  
SAP Central Services (SCS) Load Balancer and Rules  
Web Dispatcher Load Balancer and Rules  
Application, SCS, and Web Dispatcher Availability Sets  
Application VMs (based on a count)  
SCS VMs (Stand Alone or 2 VMs for High Availability)  
Web Dispatcher VMs (based on a count)



## Jumpboxes

Linux based VM Jumpbox  
Windows based VM Jumpbox  
Run Time Instance VM for Ansible Configuration  
*(Note: The Run Time Instance is a Linux jumpbox configured with the Ansible component for configuring the other Virtual machines deployed)*

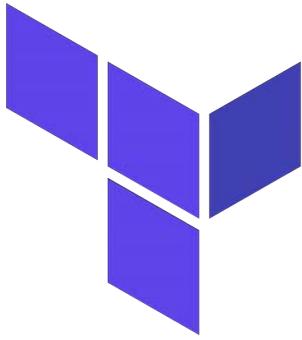
# Scenarios



# Step 1 – prepare the environment

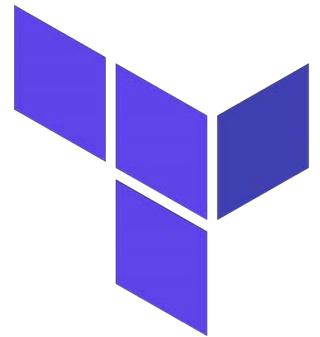
- Install Terraform
- Install Azure CLI
- Install Ansible
- Prepare your subscription: Service Principal or MSI
  - `az ad sp create-for-rbac --name <service-principal-name>`
  - Assign Contributor role for the required scope
- Clone the repo and create a Terraform Workspace
- Configure Terraform Azure RM provider

- export ARM\_SUBSCRIPTION\_ID=\${subscription\_id}
- export ARM\_TENANT\_ID=\${tenant\_id}
- export ARM\_CLIENT\_ID=\${client\_id}
- export ARM\_CLIENT\_SECRET=\${client\_secret}



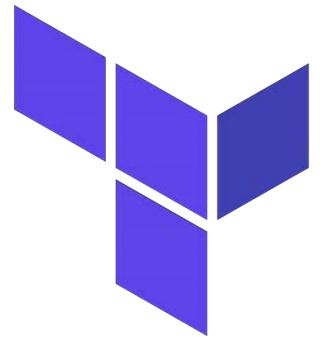
## Step 2 - Choose your scenario

rukawata Support HANA 2.00.050 on RHEL 7.6 (#837) <a href="#">...</a>	
..	
<a href="#"> application_with_ha.json</a>	add sid to application ( <a href="#">#633</a> )
<a href="#"> application_with_ha_in_existing_subnet.json</a>	add sid to application ( <a href="#">#633</a> )
<a href="#"> application_wo_ha.json</a>	add sid to application ( <a href="#">#633</a> )
<a href="#"> application_wo_ha_customimage.json</a>	add sid to application ( <a href="#">#633</a> )
<a href="#"> application_wo_ha_windows.json</a>	add sid to application ( <a href="#">#633</a> )
<a href="#"> clustered_hana.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> clustered_hana_use_custom_image_with_sbd.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> clustered_hana_with_sbd.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> clustered_hana_with_sbd_in_existing_subnet.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> ha_oracle.json</a>	Add support to AnyDB ( <a href="#">#661</a> )
<a href="#"> rti_only.json</a>	adjust default logic in app and adjust templates ( <a href="#">#595</a> )
<a href="#"> saplandscaperunner.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> single_node_hana.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> single_node_hana_single_container.json</a>	adjust default logic in app and adjust templates ( <a href="#">#595</a> )
<a href="#"> single_node_hana_with_custom_image.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> single_node_hana_with_jumpboxes.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> single_node_hana_with_xsa.json</a>	Support HANA 2.00.050 on RHEL 7.6 ( <a href="#">#837</a> )
<a href="#"> single_node_oracle.json</a>	Add support to AnyDB ( <a href="#">#661</a> )
<a href="#"> single_node_oracle_with_custom_image.json</a>	Add support to AnyDB ( <a href="#">#661</a> )



# Step 3 – Update the deployment template

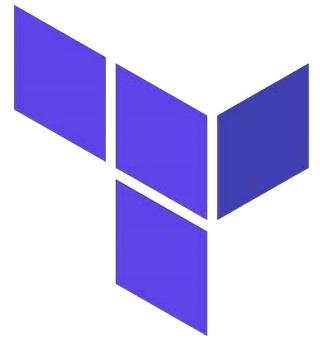
- The SAP environments deployed are configured by JSON input files. These configuration files provide a high degree of customization for the user



# Initialize Terraform

```
terraform init -var-file=../sap-  
hana/deploy/template_samples/single_node_hana.json ..../sap-  
hana/deploy/terraform
```

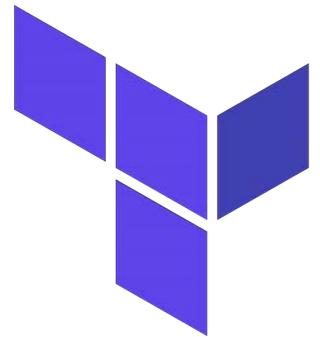
*The terraform init command is used to initialize a working directory containing Terraform configuration files. This is the first command that should be run after writing a new Terraform configuration or cloning an existing one from version control. It is safe to run this command multiple times.*



# Preview deployment

```
terraform plan -var-file=../sap-  
hana/deploy/template_samples/single_node_hana.json ..../sap-  
hana/deploy/terraform
```

*The terraform plan command is used to create an execution plan. Terraform performs a refresh, unless explicitly disabled, and then determines what actions are necessary to achieve the desired state specified in the configuration files.*

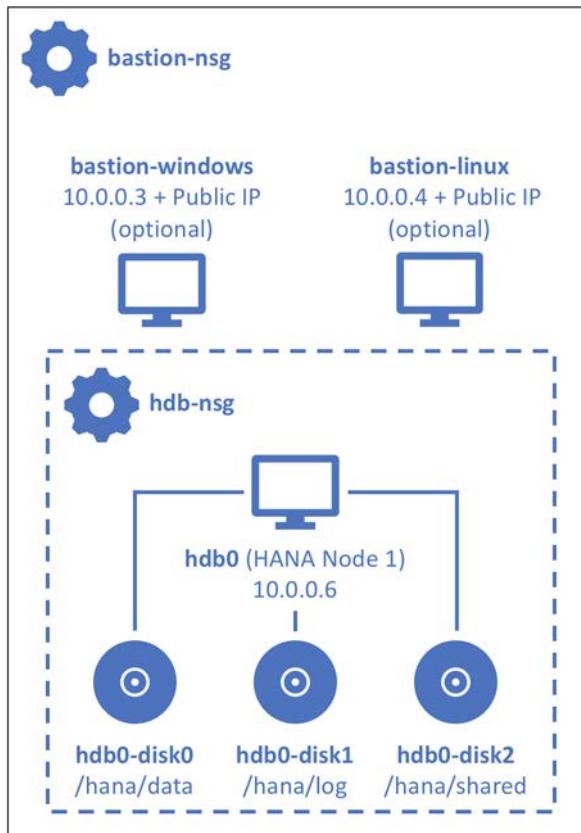
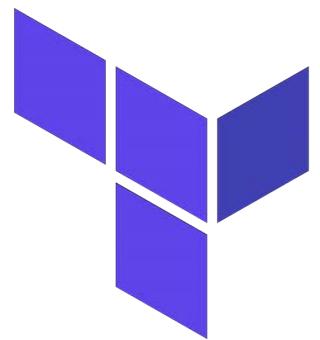


# Apply deployment

```
terraform apply -var-file=../sap-hana/deploy/template_samples/single_node_hana.json ../sap-hana/deploy/terraform
```

*The terraform apply command is used to apply the changes required to reach the desired state of the configuration, or the pre-determined set of actions generated by a terraform plan execution plan.*

# Infrastructure deployment completed



hdb1-0	Virtual machine
hdb1-0-admin-nic	Network interface
hdb1-0-backup-0	Disk
hdb1-0-data-0	Disk
hdb1-0-db-nic	Network interface
hdb1-0-log-0	Disk
hdb1-0-osdisk	Disk
hdb1-0-sap-0	Disk
hdb1-0-shared-0	Disk
HN1_hdb-alb	Load balancer
HN1_hdb-avset	Availability set
nsg-admin	Network security group
nsg-db	Network security group
nsg-mgmt	Network security group
rti	Virtual machine
rti-nic1	Network interface
rti-osdisk	Disk
rti-public-ip	Public IP address
sabootdiagc96a3f18	Storage account
sapbitsc96a3f18	Storage account
test-ppg	Proximity placement group
vnet-mgmt	Virtual network
vnet-sap	Virtual network

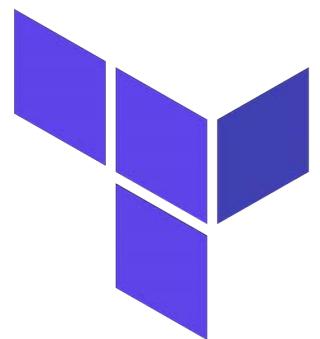
## Outputs:

```
ansible-jumpbox-public-ip-address = 52.170.212.  
ansible-jumpbox-username = azureadm
```

# Configuring SAP with Ansible



- The infrastructure has been deployed with Terraform
  - We will use Ansible to configure SAP
- Note: you can automate the entire process by specifying
- ```
"options": {  
    "enable_secure_transfer": true,  
    "ansible_execution": true,  
    "enable_prometheus": true  
},
```



# RTI + Jumpbox configuration

```
module.jumpbox.null_resource.prepare-rti (remote-exec): Selecting previously unselected package dpkg-dev.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../36-dpkg-dev_1.19.0.5ubuntu2.3_all.deb ...
Progress: [ 35%] [#####.....] ti (remote-exec): Unpacking dpkg-dev (1.19.0.5ubuntu2.3) ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Selecting previously unselected package build-essential.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../37-build-essential_12.4ubuntu1_amd64.deb ...
Progress: [ 36%] [#####.....] ti (remote-exec): Unpacking build-essential (12.4ubuntu1) ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Selecting previously unselected package python3-lib2to3.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../38-python3-lib2to3_3.6.9-1~18.04_all.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking python3-lib2to3 (3.6.9-1~18.04) ...
Progress: [ 37%] [#####.....] ti (remote-exec): Selecting previously unselected package python3-distutils.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../39-python3-distutils_3.6.9-1~18.04_all.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking python3-distutils (3.6.9-1~18.04) ...
Progress: [ 38%] [#####.....] ti (remote-exec): Selecting previously unselected package dh-python.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../40-dh-python_3.20180325ubuntu2_all.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking dh-python (3.20180325ubuntu2) ...
Progress: [ 39%] [#####.....] ti (remote-exec): Selecting previously unselected package libfakeroot:amd64.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../41-libfakeroot_1.22-2ubuntu1_amd64.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking libfakeroot:amd64 (1.22-2ubuntu1) ...
Progress: [ 40%] [#####.....] ti (remote-exec): Selecting previously unselected package fakeroot.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../42-fakeroot_1.22-2ubuntu1_amd64.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking fakeroot (1.22-2ubuntu1) ...
Progress: [ 41%] [#####.....] ti (remote-exec): Selecting previously unselected package libalgorithm-diff-perl.
module.jumpbox.null_resource.prepare-rti (remote-exec): Preparing to unpack .../43-libalgorithm-diff-perl_1.19.0-2_all.deb ...
module.jumpbox.null_resource.prepare-rti (remote-exec): Unpacking libalgorithm-diff-perl (1.19.0-2) ...
```

# Ansible playbook

- Playbooks are one of the core features of Ansible and tell Ansible what to execute. They are like a to-do list for Ansible that contains a list of tasks. Playbooks contain the steps which the user wants to execute on a particular machine. Playbooks are run sequentially.

# Connect to the Bastion to run the playbook

```
azureadm@rti:~$ ls
export-clustering-sp-details.sh hosts hosts.yml output.json sap-hana
azureadm@rti:~$ ansible-playbook -i hosts.yml ~/sap-hana/deploy/ansible/sap_playbook.yml

PLAY [localhost] ****
TASK [Gathering Facts] ****
ok: [localhost]

TASK [Include output JSON] ****
ok: [localhost]

TASK [Create dictionary with HANA database information from output.JSON] ****
ok: [localhost] => (item={'authentication': {'type': 'key', 'username': 'azureadm'}, 'components': {'hana_database': []}, 'credentials': {'cockpit_admin_password': '', 'cluster_password': '', 'os_sapadm_password': 'Help4you', 'os_sidadm_password': 'Help4you', 'xsa_admin_password': 'Manager1'}, 'db_version': '2.00.050', 'instance_number': '00', 'sid': 'HAN1'}, 'loadbalancer': {'frontend_ip': '10.1.2.4'}, 'nodes': [{'dbname': 'hdb1-0', 'ip_admin_nic': '10.1.1.10', 'ip_db_nic': '10.1.1.10', 'node_ip': '10.1.1.10', 'node_name': 'hanadbnode01', 'node_type': 'hana'}, {'dbname': 'hdb1-1', 'ip_admin_nic': '10.1.1.11', 'ip_db_nic': '10.1.1.11', 'node_ip': '10.1.1.11', 'node_name': 'hanadbnode02', 'node_type': 'hana'}, {'dbname': 'hdb1-2', 'ip_admin_nic': '10.1.1.12', 'ip_db_nic': '10.1.1.12', 'node_ip': '10.1.1.12', 'node_name': 'hanadbnode03', 'node_type': 'hana'}], 'platform': 'HANA', 'shine': {'email': 'shinedemo@microsoft.com'}, 'size': 'Demo')

PLAY [hanabdnodes] ****
TASK [Gathering Facts] ****
The authenticity of host '10.1.1.10 (10.1.1.10)' can't be established.
ECDSA key fingerprint is SHA256:SKcF1ndlPU0c+GSjhE461W1heBmTIv8NpF+YHs14YTo.
Are you sure you want to continue connecting (yes/no)? 
```

# Azure Pipelines

Cloud-hosted pipelines for Linux, Windows and macOS, with unlimited minutes for open source



## Any language, any platform, any cloud

Build, test, and deploy Node.js, Python, Java, PHP, Ruby, C/C++, .NET, Android, and iOS apps. Run in parallel on Linux, macOS, and Windows. Deploy to Azure, AWS, GCP or on-premises



## Extensible

Explore and implement a wide range of community-built build, test, and deployment tasks, along with hundreds of extensions from Slack to SonarCloud. Support for YAML, reporting and more



## Containers and Kubernetes

Easily build and push images to container registries like Docker Hub and Azure Container Registry. Deploy containers to individual hosts or Kubernetes.



## Best-in-class for open source

Ensure fast continuous integration/continuous delivery (CI/CD) pipelines for every open source project. Get unlimited build minutes for all open source projects with up to 10 free parallel jobs across Linux, macOS and Windows

Enabling feature flags for Preview Attachment and Grid Views

AdventureWorks/PackageFramework master #889

Windows Job  
Running 1m 53s

Linux Job  
Running 3m 29s

macOS Job  
Running 3m 07s

Linux Job  
Agent: Hosted Linux

- Prepare job
- Initialize job
- Get sources
- Cmdline
- Nodetool
- Install dependencies

```
yarn install v1.7.0
$ node build/npm/preinstall.js
[1/4] Resolving packages...
[2/4] Fetching packages...
[3/4] Linking dependencies...
[4/4] Building fresh packages...
$ npm run compile
#####
> code-oss-dev-build@1.0.0 compile ./adventureworks/build
> tsc -p tsconfig.build.json

✖ Done in 4.89s.
$ node ./postinstall
[#] 2/2 removed './adventureworks/extensions/node_modules/typescript/lib/tsserverlibrary.d.ts'
removed './adventureworks/extensions/node_modules/typescript/lib/tsserverlibrary.js'
removed './adventureworks/extensions/node_modules/typescript/lib/tsserverlibrary.js'
removed './adventureworks/extensions/node_modules/typescript/lib/typescriptServices.d.ts'
removed './adventureworks/extensions/node_modules/typescript/lib/typescriptServices.js'
```



<https://azure.com/pipelines>

# Azure Pipeline – Concepts

**Agents** – compute to perform build and release

- Microsoft Managed
- Self-hosted (OnPrem, Azure, Anywhere)

**Sources** : Azure repo, GitHub, Svn ...

**Jobs**: represents a series of steps that runs in sequence, multiple Jobs can run in Parallel

**Tasks**: Building block of the pipeline, task is simply a packaged script or procedure that has been abstracted with a set of inputs.

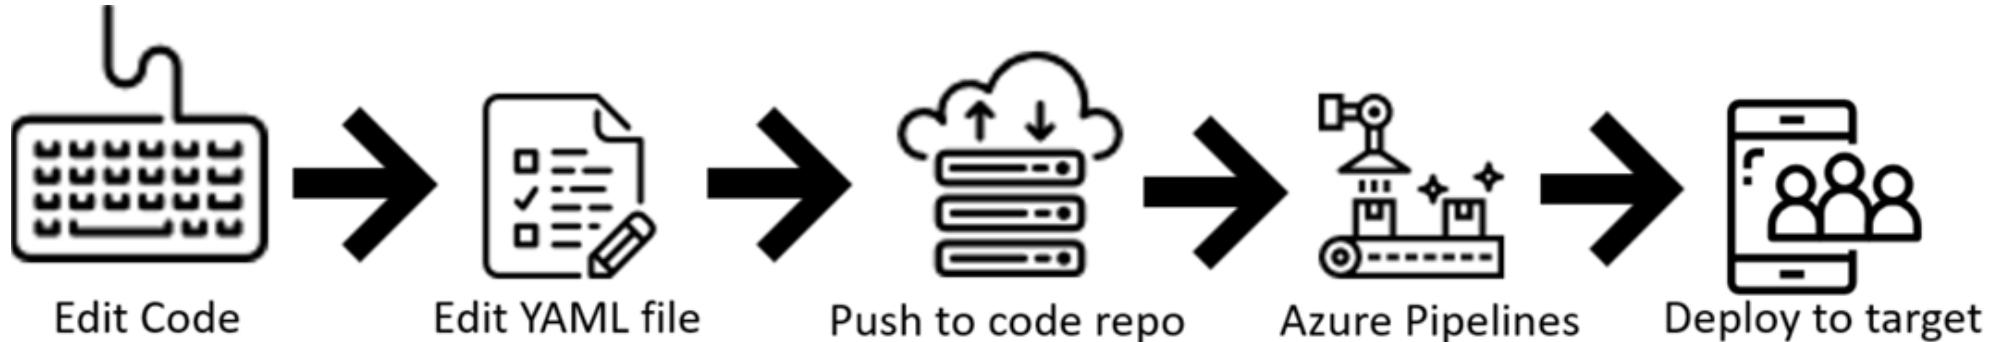
When you run a **job**, all **tasks** run in sequence on an **agent**

# Azure Pipeline - Build

- You define your Build pipeline using Agents / jobs / tasks
- Definition can be done using the UI or YAML (recommended)
- YAML can be versioned in Azure Repo
- The Build pipeline generates Artifacts, test results ...

# Azure Pipelines and YAML

- Configure your pipelines in a YAML file that exists alongside your code



- Configure Azure Pipelines to use your Git repo
- Edit your `azure-pipelines.yml` file to define your build
- Push your code to your version control repository
- Your code is now updated, built, tested, and packaged

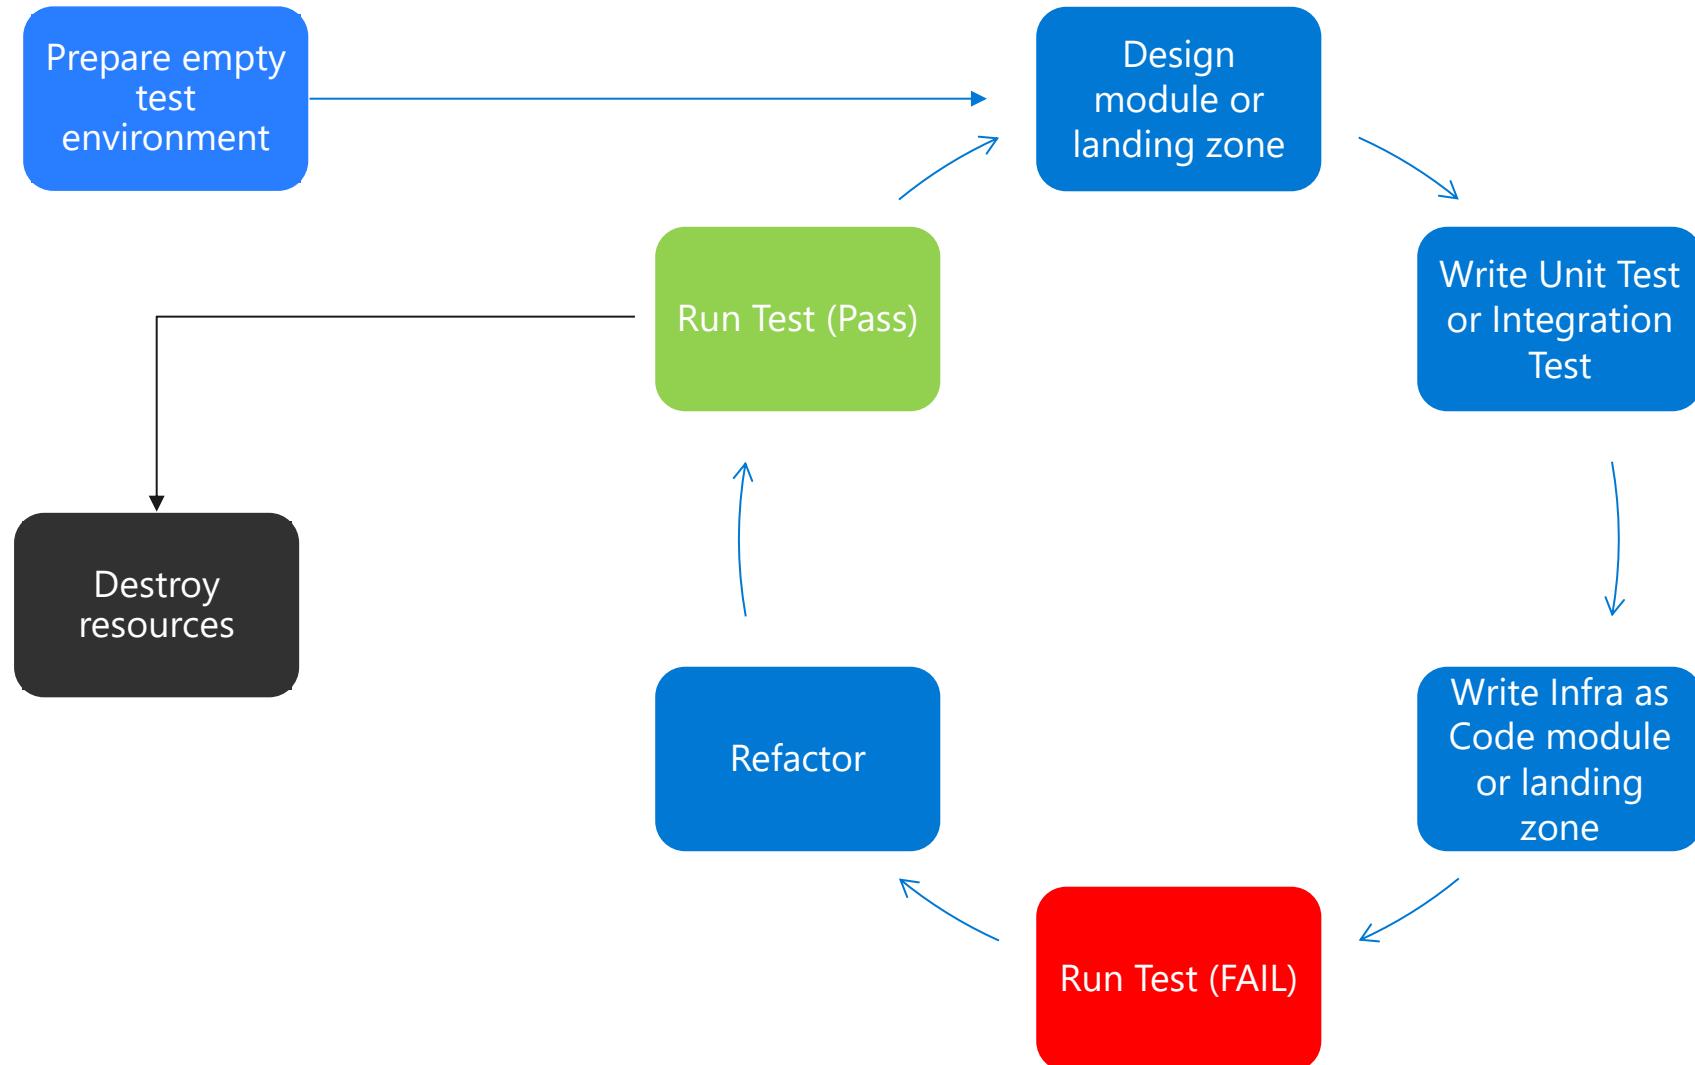
# Azure pipeline

Stages   Jobs

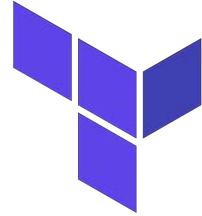
The diagram illustrates a multi-stage pipeline structure. It starts with a 'landingzone\_agw' stage containing two jobs: 'destroy landingzone\_agw' and 'destroy landingzone\_agw'. This stage has a 'Rerun stage' button. Following it is a 'landing\_zone\_ase' stage with three jobs: 'destroy landingzone\_ase', 'destroy landingzone\_ase', and 'destroy landingzone\_ase'. This stage also has a 'Rerun stage' button. Next is a 'landing\_zone\_network' stage with two jobs: 'destroy landingzone\_network' and 'destroy landingzone\_network'. This stage has a 'Rerun stage' button. Then comes a 'landing\_zone\_foundations' stage with two jobs: 'destroy landingzone\_foundations' and 'destroy landingzone\_foundations'. This stage has a 'Rerun stage' button. After that is a 'devops\_release\_aggregator' stage with five jobs: 'destroy level1 release', 'destroy level2 release', 'destroy level3 release', 'destroy level4 release', and 'destroy level0 release'. This stage has a 'Rerun stage' button. Finally, there is a 'launchpad' stage with one job: 'destroy level0 launch'. This stage has a 'Rerun stage' button. A vertical line connects the 'landingzone\_agw' stage to the 'landing\_zone\_sql' and 'landing\_zone\_apim' stages, which are grouped together.

| Sources | Repository                            | Branch / tag | Version | Related |
|---------|---------------------------------------|--------------|---------|---------|
|         | caf-configuration<br>Azure Repos      | master       | cb38360 | None    |
|         | terraform-landingzones<br>Azure Repos | master       | 7fdb582 |         |

# Testing flow with Infrastructure as Code



# Recommended tools



- `terraform show -json`: show a JSON representation of a tfstate file
- `terraform output -json`: extract the value of an output variable from the state file

`./jq`

- `jq` is a lightweight and flexible command-line JSON processor



- Azure SDK and Resource Graph API to retrieve ARM resources at scale using KQL

# Unit testing strategy

Objective: validate the configuration of a single component (e.g the terraform code for our module is working as expected)

Stage : terraform `plan` – before any deployment

Methodology :

1. Validate that the plan phase is successful
2. Retrieve the terraform `plan` state file as json
3. Parse the JSON file using your preferred tool or language
4. Extract relevant metadata and compare them to the expected result

# Integration testing strategy

Objective: validate the deployment of a module and broader landing zone is successful in the targeted environment

Stage : terraform `apply` – after the deployment in an environment.

Methodology :

1. Validate that the apply phase is successful
2. Retrieve the terraform `output` values as json (e.g resource name)
3. Extract relevant metadata and compare them to the expected result
4. Destroy must be successful (end-to-end test)

# Azure Resource Graph

Type of Azure Resource Manager object that we target  
(e.g Resources, ResourceContainers)

## Resources

```
| where type =~ "microsoft.sql/servers"  
| where name contains "svx1-sql-contososg"  
| project id, name, type, properties, tags
```

Query filtering

Select the columns to include, rename or drop,  
and insert new computed columns.

# Azure Resource Graph – integration testing

Written in GoLang and leverages the Azure SDK

`go -t ./...`

1. Uses values from terraform output (as env variables)
2. Calls the Azure Resource Graph query
3. Validates the number of results
4. Optional parse the retrieved data (tags, properties)

```
func TestSQLServer(t *testing.T) {
    t.Parallel()

    serverName := os.Getenv("SQL_SERVER_NAME")
    rgName := os.Getenv("RG_NAME")
    location := os.Getenv("LOCATION")
    kind := os.Getenv("SQL_VERSION")
    env := os.Getenv("ENVIRONMENT")

    var query string = fmt.Sprintf("Resources | where type =~ 'microsoft.sql/servers' |"
        "subscriptionID := os.Getenv("ARM_SUBSCRIPTION_ID")
        fmt.Println("getting client")
        fmt.Println(query)

    graphClient := resourcegraph.New()
    authorizer, err := auth.NewAuthorizerFromEnvironment()
    if err == nil {
        graphClient.Authorizer = authorizer
    }

    queryRequest := new(resourcegraph.QueryRequest)
    queryRequest.Query = &query
    queryRequest.Subscriptions = &([]string{subscriptionID})
    result, err := graphClient.Resources(context.Background(), *queryRequest)
    if err != nil {
        t.Errorf("%v\n", err)
    }
    // fmt.Printf(result.Data.(map[string]interface{}).location)

    assert.Equal(t, int64(1), *result.Count)
}
```

# Session Takeaway

---

Favor Infra As Code to deploy complex architectures

Automate the deployment process and separate code from configuration (12 factor app)

Implement a DevOps pipeline

Test your SAP deployment

# Q&A

Reach out to the team  
[sap-on-azure-pe-apac@microsoft.com](mailto:sap-on-azure-pe-apac@microsoft.com)

# Feedback

Your feedback is very important for us.

<https://aka.ms/SAPAPAC-POE-FEEDBACK>





# SAP on Azure Enablement

Next Session – Migration and Architecture best practices

Next week - Monday, Oct 26, 2020, 10am SGT

Reach out to the team  
[sap-on-azure-pe-apac@microsoft.com](mailto:sap-on-azure-pe-apac@microsoft.com)



# Reach out to the team



Ravi Gangampalli



Sajit Nair



Nicolas Yuen



Inseob Kim

[sap-on-azure-pe-apac@microsoft.com](mailto:sap-on-azure-pe-apac@microsoft.com)